Project 2.2

Experiment with Bayes Classifiers

AOED Classifier:

Averaged One-Dependence Estimators (AODE) is a probabilistic classification learning technique. It was developed to address the attribute-independence problem of the popular Naive Bayes classifier. It frequently develops substantially more accurate classifiers than naive Bayes at the cost of a modest increase in the amount of computation. Like naive Bayes, AODE does not perform model selection and hence has low variance. It supports incremental learning whereby the classifier can be updated efficiently with information from new examples as they become available. It predicts class probabilities rather than simply predicting a single class, allowing the user to determine the confidence with which each classification can be made. Its probabilistic model can directly handle situations where some data are missing.

As you can see the classification error has been %50, which is noticeable.

I was very interested to see the role of humidity in our classifier. So I visualized the error vs. the humidity. As you can see, people generally want to play when the humidity is normal (and hence the error is small), but when the humidity is high, then the result depends on other factors (like wind).

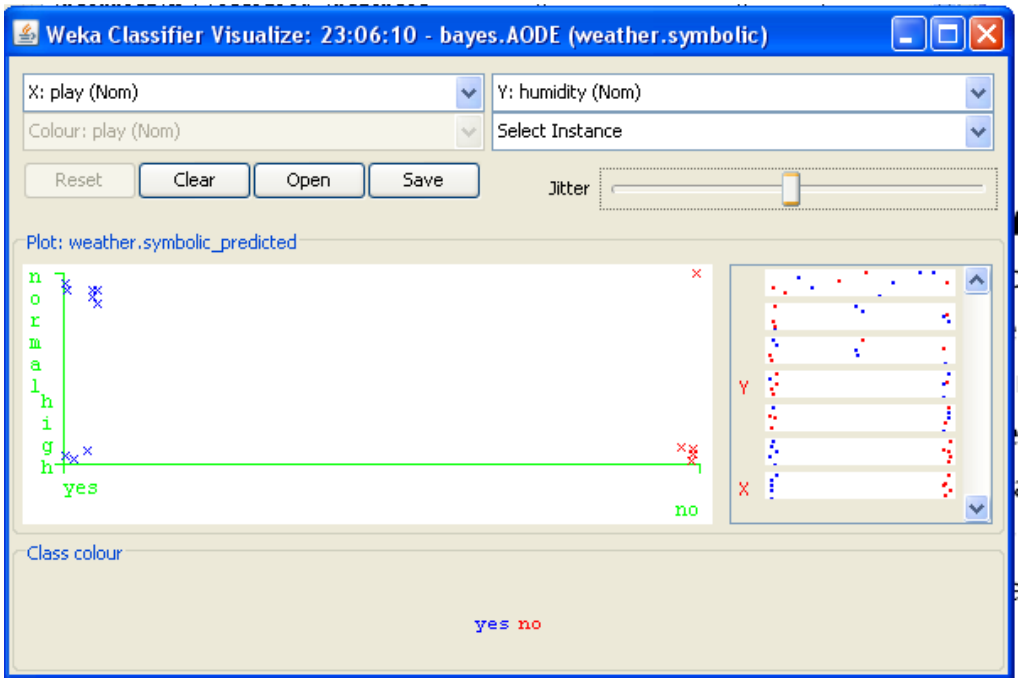

=== Run information ===

Scheme: weka.classifiers.bayes.AODE -F 1

Relation: weather.symbolic

Instances: 14

Attributes: 5

outlook

temperature

humidity

windy

play

Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

The AODE Classifier

Class yes: Prior probability = 0.63

Class no: Prior probability = 0.38

Dataset: weather.symbolic

Instances: 14

Attributes: 5

Frequency limit for superParents: 1

Correction: laplace

Time taken to build model: 0 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	*7*	*50 %*
Incorrectly Classified Instances	*7*	*50 %*
Kappa statistic	*-0.0426*	
Mean absolute error	*0.4706*	
Root mean squared error	*0.498*	
Relative absolute error	*98.8198 %*	
Root relative squared error	*100.941 %*	
Total Number of Instances	*14*	

=== Detailed Accuracy By Class ===

	TP Rate	*FP Rate*	*Precision*	*Recall*	*F-Measure*	*ROC Area*	*Class*
	0.556	*0.6*	*0.625*	*0.556*	*0.588*	*0.533*	*yes*
	0.4	*0.444*	*0.333*	*0.4*	*0.364*	*0.533*	*no*
Weighted Avg.	*0.5*	*0.544*	*0.521*	*0.5*	*0.508*	*0.533*	

=== Confusion Matrix ===

a	*b*	*<-- classified as*
5	*4*	*| a = yes*
3	*2*	*| b = no*

BayesNet

Bayes Network learning using various search algorithms and quality measures.

Base class for a Bayes Network classifier. Provides datastructures (network structure, conditional probability distributions, etc.) and facilities common to Bayes Network learning algorithms like K2 and B.Application of BayesNet gives ~57% of accuracy. Below is the run information:

=== Run information ===

Scheme: weka.classifiers.bayes.BayesNet -D -Q weka.classifiers.bayes.net.search.local.K2 -- -P 1 -S BAYES -E weka.classifiers.bayes.net.estimate.SimpleEstimator -- -A 0.5

Relation: weather.symbolic

Instances: 14

Attributes: 5

outlook

temperature

humidity

windy

play

Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

Bayes Network Classifier

not using ADTree

#attributes=5 #classindex=4

Network structure (nodes followed by parents)

outlook(3): play

temperature(3): play

humidity(2): play

windy(2): play

play(2):

LogScore Bayes: -69.07317135664013

LogScore BDeu: -83.46880542273105

LogScore MDL: -82.71568504897063

LogScore ENTROPY: -65.56181240647145

LogScore AIC: -78.56181240647145

Time taken to build model: 0 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	*8*	*57.1429 %*
Incorrectly Classified Instances	*6*	*42.8571 %*
Kappa statistic	*-0.0244*	
Mean absolute error	*0.415*	
Root mean squared error	*0.4909*	
Relative absolute error	*87.1501 %*	
Root relative squared error	*99.5104 %*	
Total Number of Instances	*14*	

=== Detailed Accuracy By Class ===

TP Rate	*FP Rate*	*Precision*	*Recall*	*F-Measure*	*ROC Area*	*Class*

0.778 0.8 0.636 0.778 0.7 0.622 yes

0.2 0.222 0.333 0.2 0.25 0.622 no

Weighted Avg. 0.571 0.594 0.528 0.571 0.539 0.622

=== Confusion Matrix ===

a b <-- classified as

7 2 | a = yes

4 1 | b = no

Naïve Bayes

This classifier is expected to give a worse accuracy because of the assumption of independence. Class for a Naive Bayes classifier using estimator classes. Numeric estimator precision values are chosen based on analysis of the training data. For this reason, the classifier is not an UpdateableClassifier (which in typical usage are initialized with zero training instances) -- if you need the UpdateableClassifier functionality, use the NaiveBayesUpdateable classifier. The NaiveBayesUpdateable classifier will use a default precision of 0.1 for numeric attributes when buildClassifier is called with zero training instances.

=== Run information ===

Scheme: weka.classifiers.bayes.NaiveBayes

Relation: weather.symbolic

Instances: 14

Attributes: 5

outlook

temperature

humidity

windy

play

Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

Naive Bayes Classifier

Class

Attribute yes no

(0.63) (0.38)

==========

outlook

sunny 3.0 4.0

overcast 5.0 1.0

rainy 4.0 3.0

[total] 12.0 8.0

temperature

hot 3.0 3.0

mild 5.0 3.0

cool 4.0 2.0

[total] 12.0 8.0

humidity

high 4.0 5.0

normal 7.0 2.0

[total] 11.0 7.0

windy

TRUE 4.0 4.0

FALSE 7.0 3.0

[total] 11.0 7.0

Time taken to build model: 0 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances 8 57.1429 %

Incorrectly Classified Instances 6 42.8571 %

Kappa statistic -0.0244

Mean absolute error 0.4374

Root mean squared error 0.4916

Relative absolute error 91.8631 %

Root relative squared error 99.6492 %

Total Number of Instances 14

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	ROC Area	Class
	0.778	0.8	0.636	0.778	0.7	0.578	yes
	0.2	0.222	0.333	0.2	0.25	0.578	no
Weighted Avg.	0.571	0.594	0.528	0.571	0.539	0.578	

=== Confusion Matrix ===

a	b	<-- classified as
7	2	a = yes
4	1	b = no

Experiment with ANN

Multi-layer Perceptron

Experiment with ANN gives ~75 accuracy at best with variations of number of hidden layers.

=== Run information ===

Scheme: *weka.classifiers.functions.MultilayerPerceptron -L 0.3 -M 0.2 -N 500 -V 0 -S 0 -E 20 -H a*

Relation: *weather.symbolic*

Instances: *14*

Attributes: *5*

outlook

temperature

humidity

windy

play

Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

Sigmoid Node 0

Inputs Weights

Threshold -4.597967080790813

Node 2 2.433270074007239

Node 3 2.0546443732203774

Node 4 1.364159803860347

Node 5 2.6974766889493536

Node 6 3.908322709064356

Sigmoid Node 1

Inputs Weights

Threshold 4.601251960011152

Node 2 -2.4045226373071156

Node 3 -2.0532744956144127

Node 4 -1.379986429753948

Node 5 -2.756274547604192

Node 6 -3.877948258791871

Sigmoid Node 2

Inputs Weights

Threshold -0.1550798021501342

Attrib outlook=sunny -1.323464477913686

Attrib outlook=overcast 1.6602675280399888

Attrib outlook=rainy -0.3207802552865604

Attrib temperature=hot -0.2873122456981835

Attrib temperature=mild 1.181190360097958

Attrib temperature=cool -0.7853150475848826

Attrib humidity 2.808930687905

Attrib windy 1.9190213581350706

Sigmoid Node 3

Inputs Weights

Threshold -0.18031675012278034

Attrib outlook=sunny -1.1524514010228344

Attrib outlook=overcast 1.5760227701429683

Attrib outlook=rainy -0.32578400279223824

Attrib temperature=hot -0.2760307631136823

Attrib temperature=mild 1.0450876279343007

Attrib temperature=cool -0.6318819517738498

Attrib humidity 2.4504774603875408

Attrib windy 1.678251292646871

Sigmoid Node 4

Inputs Weights

Threshold -0.3554146745674961

Attrib outlook=sunny -0.46574052680925143

Attrib outlook=overcast 1.4382073898080827

Attrib outlook=rainy -0.6194183985830608

Attrib temperature=hot -0.0670794406887232

Attrib temperature=mild 0.6337484752708613

Attrib temperature=cool -0.20814280117719502

Attrib humidity 1.982466584793048

Attrib windy 0.9946423645131915

Sigmoid Node 5

Inputs Weights

Threshold -0.06888405078498452

Attrib outlook=sunny -1.3982064219096493

Attrib outlook=overcast 1.8084944112736516

Attrib outlook=rainy -0.31997269602762973

Attrib temperature=hot -0.3035821635771427

Attrib temperature=mild 1.2908528760310662

Attrib temperature=cool -0.8921466424329777

Attrib humidity 3.1090049574873424

Attrib windy 2.0747113212966872

Sigmoid Node 6

Inputs Weights

Threshold 0.04399369934901554

Attrib outlook=sunny -1.80182134279014

Attrib outlook=overcast 2.2544547024444554

Attrib outlook=rainy -0.40095717506501327

Attrib temperature=hot -0.41558677311306397

Attrib temperature=mild 1.589170285947685

Attrib temperature=cool -1.2545441906677217

Attrib humidity 4.119310666164331

Attrib windy 2.740851006387263

Class yes

Input

Node 0

Class no

Input

Node 1

Time taken to build model: 0.45 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	*10*	*71.4286 %*
Incorrectly Classified Instances	*4*	*28.5714 %*
Kappa statistic	*0.3778*	
Mean absolute error	*0.287*	
Root mean squared error	*0.5268*	
Relative absolute error	*60.2616 %*	
Root relative squared error	*106.7798 %*	
Total Number of Instances	*14*	

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	ROC Area	Class
	0.778	0.4	0.778	0.778	0.778	0.778	yes
	0.6	0.222	0.6	0.6	0.6	0.778	no
Weighted Avg.	0.714	0.337	0.714	0.714	0.714	0.778	

=== Confusion Matrix ===

a	b	<-- classified as
7	2	a = yes
2	3	b = no

Voted Perceptron

This network works much worse and gives ~57% accuracy at best. Below are the run information.

=== Run information ===

Scheme: *weka.classifiers.functions.VotedPerceptron -I 1 -E 1.0 -S 1 -M 10000*

Relation: *weather.symbolic*

Instances: *14*

Attributes: *5*

outlook

temperature

humidity

windy

play

Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

VotedPerceptron: Number of perceptrons=4

Time taken to build model: 0 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances	*8*	*57.1429 %*
Incorrectly Classified Instances	*6*	*42.8571 %*
Kappa statistic	*-0.1351*	
Mean absolute error	*0.3918*	
Root mean squared error	*0.5825*	
Relative absolute error	*82.2829 %*	
Root relative squared error	*118.0664 %*	
Total Number of Instances	*14*	

=== Detailed Accuracy By Class ===

TP Rate FP Rate Precision Recall F-Measure ROC Area Class

0.889 1 0.615 0.889 0.727 0.622 yes

0 0.111 0 0 0 0.622 no

Weighted Avg. 0.571 0.683 0.396 0.571 0.468 0.622

=== Confusion Matrix ===

a b <-- classified as

8 1 | a = yes

5 0 | b = no

Winnow Network

Winnow network works based on sparse connections and hence can't beat the multi-layer perceptron. Here it gives only 50% of accuracy. Below is the run information.

=== Run information ===

Scheme: weka.classifiers.functions.Winnow -I 1 -A 2.0 -B 0.5 -H -1.0 -W 2.0 -S 1

Relation: weather.symbolic

Instances: 14

Attributes: 5

outlook

temperature

humidity

windy

play

Test mode: 10-fold cross-validation

=== Classifier model (full training set) ===

Winnow

Attribute weights

w0 8.0

w1 1.0

w2 2.0

w3 4.0

w4 2.0

w5 2.0

w6 1.0

w7 1.0

Cumulated mistake count: 7

Time taken to build model: 0 seconds

=== Stratified cross-validation ===

=== Summary ===

Correctly Classified Instances 7 50 %

Incorrectly Classified Instances 7 50 %

Kappa statistic -0.2564

Mean absolute error 0.5

Root mean squared error 0.7071

Relative absolute error 105 %

Root relative squared error 143.3236 %

Total Number of Instances 14

=== Detailed Accuracy By Class ===

	TP Rate	FP Rate	Precision	Recall	F-Measure	ROC Area	Class
	0.778	1	0.583	0.778	0.667	0.389	yes
	0	0.222	0	0	0	0.389	no
Weighted Avg.	0.5	0.722	0.375	0.5	0.429	0.389	

=== Confusion Matrix ===

a	b	<-- classified as
7	2	a = yes
5	0	b = no

Discuss Your Results

As mentioned before, due to dependence on sparse connections, Winnow is the worst choice for the network when we have few attributes and a small database. Winnow is good when we are dealing with numerous attributes and large dataset which contain redundancy. With our dataset, Multi-layer Perceptron gives the best results.